


# HW 3: MDP, Gridworld, Bellman's Equation


Nicholas Strohmeyer (nas3882)

## Part 1: Value Iteration

- Various heatmaps allow us to visualize value distributions given different end goal states
- In all cases value emanates out from the center target state (where the reward is densest)
- State space sweeps to converge: **54** (case 1: when goal = (19,9))




Below is the policy generated by value iteration in case 1 where the goal = (19,9). In part 2, we will look



at the policy function using policy iteration and some smoother/ cleaner lines will emerge. However, this policy is generally correct. If we are to the right, the policy says go left (sky blue), if we are to the left, the policy says go right (white). For most of the area under the goal state, the policy prescribes an up action. We will see that sometimes it will be slightly more efficient to go right or left (depending on the state) when we look at the policy iteration map. The silver line across the top indicating "down" is interesting. I explain why I believe the algorithms prescribed these actions in part 2 (it has to do with

the boundary). Lastly, I include a plot of the prescribed trajectory from 2 different starting states: (0,0) and (0,19):

**Steps taken = 28**

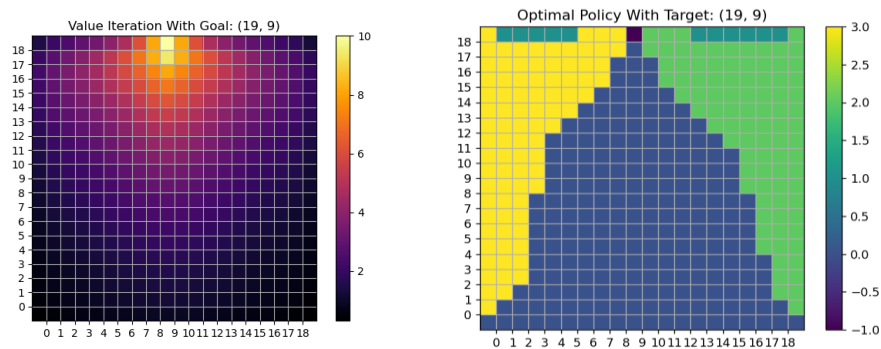


## Part 2: Policy Iteration

- State Space Sweeps when goal (19,9): **49**

We see policy iteration did better (i.e. required less sweeps of the state space to converge) We could likely converge even faster with a better initial guess for the policy (right now it is randomized). For example, we could instead try an initial policy such as: “go right if column < 9, go left if column > 9, and go up if column = 9” (based on our knowledge of starting state and goal state)

- I come back to test this hypothesis in part 3



### **Left Map:**

The inferno map on left = optimal value function. I arbitrarily made the target state a much higher reward than surrounding states in order to help convergence along (that is why the scale goes up to 10). We can see that value emanates outward from grid square (19,9) almost like a vector field pulling surrounding particles in towards the target state. We can see this function is essentially identical to the value iteration algorithm result.

### **Right Map:**

4-toned map represents 4 actions (the reward state is marked -1, colored purple to indicate no action taken once reached) 0 = up (blue), 1 = down (teal), 2 = left (green), 3 = right (yellow)

Generally the optimal policy is doing as we would expect. If we are directly below the target state, the optimal action is to move up. If we are further off to the right, the optimal action is to go left and vice versa.

### **Why do we move down along the top row ?**

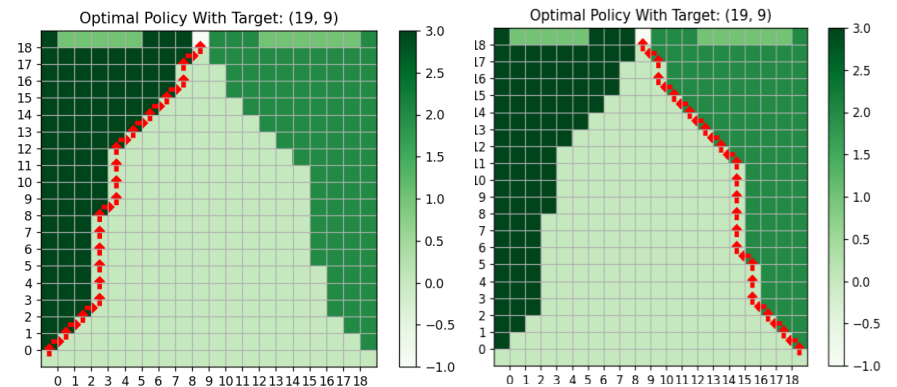
The interesting exception to this rule is across the top row where our target state lives (that is  $y = 19$ ). Unless we are right next to (19,9), the optimal policy actually tells us to go down. The reason for this is not intuitive at first, but after thinking about it for a little while, I think it makes sense. It is the effect of calculating value along the boundary where one of the actions is illegal (we cannot go up). These states end up with one less grid square contributing to the overall expectation in bellman's equation. The effect is that states immediately below (19,9) have slightly higher reward values than those to the immediate left and right. This effect emanates to the left and right (in row  $y = 18$ ) and thus gives the

down action a slightly higher reward in those aforementioned squares (19, 1-5) & (19, 12-18). Fortunately, closer to (19,9) the policy still correctly identifies the optimal move to the left and right.

Using the optimal policy function, I simulated a trajectory starting from 2 states again: (0,0) and (19,0).

It is interesting to see how this moves exactly along the boundary of the “up” and “right” regions. This makes sense. Deviating too deeply into either region would represent inefficiency (ie wasted movements).


This is a satisfying visual representation of the optimality of bellman’s equation. The agent is not only finding the target state but doing so in as few of steps as possible. Note: this explanation is also true of the value iteration trajectories we saw in part 1. In part 3, I consider the differing geometries between the optimal policy functions of the 2 algorithms




**Steps taken = 28**

Just for fun, here a few other simulated trajectories

Policy iteration



Value Iteration



## Part 3: Additional Considerations

### 1. Convergence Times

In part 2, I hypothesized that using a more informed initial guess would make policy iteration converge faster. But after switching to the biased policy shown on the right, and yet did not observe faster convergence. I still think it is reasonable that this would be true in larger state spaces (although did not test that here)

On the other hand, I found that changing the target state did highlight a more significant difference between policy iteration and value iteration and affected speed of convergence greatly. Specifically, value iteration consistently requires about 50-55 state space sweeps in order to converge, regardless of the goal state. However, the closer the goal state is to our initial, policy iteration can converge significantly faster. The table to the right displays this data.

Goal	Value Iteration	Policy Iteration	Delta
19,9	54	49	5
10,9	47	29	18
19,0	53	1	52
5,6	52	22	30
19,19	53	55	-2
0,19	54	49	5
1,0	55	21	34
11,1	54	38	16

```
def initPolicy(self, states, initType = ["random", "biased"]):
    if initType == "random":
        for i in range(self.size):
            for j in range(self.size):
                n = np.random.random()
                if n <= .25:
                    action = 0
                elif .25 < n <= .5:
                    action = 1
                elif .5 < n <= .75:
                    action = 2
                else:
                    action = 3
                states[i][j] = action
    else:
        for i in range(self.size):
            for j in range(self.size):
                if (j == self.goal[1]) & (i <= self.goal[0]):
                    states[i][j] = 0 # up
                elif (j == self.goal[1]) & (i > self.goal[0]):
                    states[i][j] = 1 # down
                elif j > self.goal[1]:
                    states[i][j] = 2 # left
                elif j < self.goal[1]:
                    states[i][j] = 3 # right
                else:
                    pass
    return states
```


### 2. Policies with different Goal States

Goal = (10,9) i.e. center-grid target state

I was curious to see what would happen and how the 2 algorithms would compare when we placed the goal state exactly at the center of the grid (or close to exact center since even # of squares). What type of policy would each algorithm come up with? We saw that policy iteration created slightly smoother boundaries before, but if the target were in a more “balanced” position, would the 2 policies differ so greatly?

The figures on the right show that policy iteration does indeed create cleaner boundaries again (above) as value iteration (below) is again a more “choppy” policy. This is evidenced by the rotational symmetry in the top diagram. No matter where we are in gridworld, we always tend to funnel down towards the center along a sort of “uniform vector field” on the diagonals if you will. However, value iteration is optimal as well although not as perfectly symmetric in its action space. No matter what region we are in, we always take the minimum number of steps to get to the center (that is if we need to go right, we will never go left afterwards and if we need to go up, we will never go down etc.)

My hypothesis for why the policy iteration comes up with a “smoother” action space is that it could be due to the symmetry that was built into the initial guess. It would be interesting to explore further



## Part 4: Project Code

I created 2 python files to simulate gridworld: A gridworld **class** that contains about half of the methods and the parameters for a particular construct of gridworld (ie size, initial conditions etc.) and **runGW.py**, where I put the main() routine, the plotting functions **as well as the actual policy and value iteration algorithms** (although policy/value iteration both use methods that are contained in the gridworld class). It would make more sense to keep policy/value iteration algorithms in the same file as the methods they depend on, so hopefully this is not too confusing. I added method descriptions/documentation to hopefully make a little clearer/ easier to read

### 1<sup>st</sup> File

#### runGW.py:

- Globals : allows us to set some attributes about gridworld (I use what the assignment calls for)
- plotvalueMap()
- plotPolicyMap()
- plotSimPath() - plots a simulated trajectory through gridworld
- main() -- gives user choice of running policy or value iteration – computes path starting from initial (0,0) given either case ... based on the computed value function

```
1 from gridworld import gridworld
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 # Init World/ Globals
6 setgoal = (19,9)
7
8 world = gridworld(goal = setgoal)
9 size = world.size
10 initial = world.start
11 goal = world.goal
12
13 # Value Iteration Algorithm
14 def valueIteration():
15     curr = [[0 for i in range(20)] for j in range(20)]
16     prev = [[0 for i in range(20)] for j in range(20)]
17     iters , i= 0, 0
18     while iters < 1000:
19         # calculate values for all states
20         for i in range(size):
21             for j in range(size):
22                 world.state = (i,j) # increment state
23                 val = max(world.computeValues(prev)) # update value , make sure this didnt break
24                 curr[i][j] = round(val,2)
25                 #print(curr[i][j])
26             # check convergence
27             if iters != 1:
28                 delta = world.checkStop(curr, prev)
29                 # condition is to be small 10x in a row?
30                 if(delta < .001):
```

```

31         print("iterations to converge:", iters)
32         break
33     # update prevValues to currValues
34     for i in range(size):
35         for j in range(size):
36             prev[i][j] = curr[i][j]
37     iters += 1
38     return curr
39
40 # policy evaluation , subroutine of policy iteration
41 def policyEvaluation(pi, currVal):
42     delta, iters = 0, 0
43     while (delta < .01) and iters < 1000:
44         for i in range(0,size):
45             for j in range(0,size):
46                 world.state = (i,j)
47                 a = pi[i][j]
48                 n = world.getNeighbors()
49                 v = currVal[i][j]
50                 currVal[i][j] = world.computePolicyStep(move = a, values = currVal, nbrs = n)
51                 delta = max(delta, abs(currVal[i][j] - v))
52             iters += 1
53     return currVal
54
55 # Policy Iteration Algorithm:
56 def policyIteration():
57     #V = value function to converge to optimal, Pi to converge on best policy
58     empty = [[0 for i in range(20)] for j in range(20)]
59     Pi = world.initPolicy(empty)
60     V = [[0 for i in range(20)] for j in range(20)]

```

```

61     V = policyEvaluation(Pi, V) # init eval before improvement
62     stable = False
63
64     # improvement step:
65     n = 0
66     while (stable == False) and n < 100:
67         stable = True
68         for i in range(20):
69             for j in range(20):
70                 world.state = (i,j)
71                 if (i,j) == world.goal:
72                     Pi[i][j] = -1
73                     continue
74                 best = np.argmax(world.computeValues(V))
75                 mxm = max(world.computeValues(V))
76                 if best != Pi[i][j]:
77                     stable = False
78                     Pi[i][j] = best
79             if stable == False:
80                 V = policyEvaluation(Pi, V)
81         else:
82             print("State Space Sweeps: ", n)
83             n += 1
84     return Pi, V
85

```

- Skipping past some of the plotting functions

```

def main():
    #value iteration
    #optValue = valueIteration()

    # policy iteration
    optPolicy, optValue = policyIteration()
    path, moves = world.simulateWorld(Pi = optPolicy)
    -----

    # Plots
    plotValueMap(optValue) # Value Function:
    # Policy plot
    fig = plt.figure()
    ax = fig.add_subplot(1, 1, 1)
    ticks = np.arange(.5, 19.5, 1)
    tick_labels = np.arange(0,19,1)
    ax.set_xticks(ticks, labels = tick_labels)
    ax.set_yticks(ticks, labels = tick_labels)
    plotPolicyMap(optPolicy)
    plt.grid()
    #plt.legend(labels = "0: up 1: down 2: left 3: right")
    plt.colorbar()
    #plotSimPath(path, optPolicy)

    plt.show()
    return

main()

```

2<sup>nd</sup> File:

gridworld.py (class)

- **`_init_`** -- constructor
- **`move()`** -- given an action input, update the state of gridworld. If action is invalid (i.e boundary) do nothing
- **`initPolicy()`** – randomizes a starting policy configuration for the policy iteration algorithm (could use an initial guess to get better performance) Allows a choice between a random initial policy or directed one (see code above in part 3 for “biased” version)
- **`extractPolicy()`** – using only an optimal value function as input, creates the optimal policy function by taking argmax of the nearest neighbors in each state
- **`computeValues()`** – computes bellman’s equation for all possible actions at a given state (returns all values and then we take argmax() in outer routine)
- **`computeStep()`** -- computes bellman’s equation for a single constrained policy at a given state
- **`checkStop()`** -- checks a convergence criterion for value iteration algorithm
- **`simulateWorld()`** – given a policy function and initial state, simulates a trajectory. Built-in stopping condition in case the trajectory is periodic (infinite loop)
- **`getNeighbors()`** -- given a state, find the coordinates of each of its neighbors. If adjacent to boundary in some direction, returns None for that direction

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3  import random
4
5  class gridworld(object):
6      def __init__(self, size = 20, start = (0,0), goal = (19,9), p = 0.6, gamma = .95):
7          self.size, self.goal, self.p, self.gamma = size, goal, p, gamma
8          self.start, self.state = start, start
9          self.optValues = np.zeros((self.size, self.size))
10         self.optPolicy = np.zeros((self.size, self.size))
11
12         # [0,1,2,3] = [up, down, right, left]
13         def move(self, action):
14             if action == 0: # up
15                 if self.state[0] < self.size - 1:
16                     self.state = (self.state[0] + 1, self.state[1])
17                 else:
18                     print("boundary, no movement")
19                     print(self.state)
20             elif action == 1: #down
21                 if self.state[1] > 0:
22                     self.state = (self.state[0] - 1, self.state[1])
23                 else:
24                     print("boundary, no movement")
25                     print(self.state)
26             elif action == 2: #left
27                 if self.state[1] > 0:
28                     self.state = (self.state[0], self.state[1] - 1)
29                 else:
30                     print("boundary, no movement")

```

```

31         |         print(self.state)
32         elif action == 3: #right
33             if self.state[1] < self.size - 1:
34                 self.state = (self.state[0], self.state[1] + 1)
35             else:
36                 print("boundary, no movement")
37                 print(self.state)
38         else:
39             pass
40
41         # [0,1,2,3] = [up, down, left, right]
42         def initPolicy(self, states):
43             for i in range(self.size):
44                 for j in range(self.size):
45                     n = np.random.random()
46                     if n <= .25:
47                         action = 0
48                     elif .25 < n <= .5:
49                         action = 1
50                     elif .5 < n <= .75:
51                         action = 2
52                     else:
53                         action = 3
54                     states[i][j] = action
55             return states
56
57         # neighbors: (up, down, right, left)
58         def computeValues(self, values):
59             nbrs, r = self.getNeighbors(), 0
60             # base case

```



```

61     if self.state == self.goal:
62         return 1
63     test = [0,0,0,0] # test each direction
64     for move in [0, 1, 2, 3]: # [up, down, "right", "left" ]
65         if nbrs[move] == None:
66             continue
67         else:
68             for s in nbrs:
69                 if s == None: #boundary
70                     continue
71                 if nbrs.index(s) == move: # use correct transition factor
72                     T = self.p
73                 else:
74                     T = (1-self.p)/3
75                 test[move] += T*(r+self.gamma*values[s[0]][s[1]]) # bellman's equation, r is always 0 here
76 # having trouble with direct update, np array will not set to a variable
77 # self.currValues[self.state] = max(test)
78     return test
79
80 def computePolicyStep(self, move = 0, values = [[]], nbrs = []):
81     test = 0
82     if self.state == self.goal:
83         return 10 # for some reason I needed to hardcode in a much higher reward here to converge properly
84     for s in nbrs:
85         if s == None:
86             continue
87         if s == self.goal:
88             r = 1
89         else:
90             r = 0

```

```

91         if nbrs.index(s) == move:
92             T = self.p
93         else:
94             T = (1-self.p)/3
95         test += T*(r+self.gamma*values[s[0]][s[1]])
96     return test
97
98 def checkStop(self, prev, curr):
99     diff = 0
100     for i in range(self.size):
101         for j in range(self.size):
102             diff += abs((prev[i][j] - curr[i][j]))
103     return diff
104
105 def getNeighbors(self): # check bounds and get neighbors if exists, "up" actually decreases row num
106     if self.state[0] < self.size-1:
107         up = (self.state[0] + 1, self.state[1])
108     else:
109         up = None
110     if self.state[0] > 0:
111         down = (self.state[0] - 1, self.state[1])
112     else:
113         down = None
114     if self.state[1] < self.size-1:
115         right = (self.state[0], self.state[1] + 1)
116     else:
117         right = None
118     if self.state[1] > 0:
119         left = (self.state[0], self.state[1] - 1)
120     else:

```

```

120         else:
121             left = None
122
123         return [up, down, left, right]
124
125     # based on a computed optimal policy pi and initial state, simulate the MDP
126     def simulateWorld(self, start = (0,0), Pi = [[]]):
127         self.state = start
128         states, actions = [start], []
129         i, j = start[0], start[1]
130         go, iter = True, 0
131         while (go == True) & (iter < 1000):
132             action = Pi[i][j]
133             actions.append(action)
134             self.move(action)
135             new = self.state
136             if new == self.goal:
137                 print("Target State Achieved")
138                 go = False
139             states.append(new)
140             i, j = new[0], new[1]
141             iter += 1
142         return states, actions
143
144
145

```

(added the below function after the others)

```

# generate a policy given only an optimal value function/somewhat redundant
def extractPolicy(self, V = [[]]):
    policy = [[ 0 for i in range(20)] for j in range(20)]
    for i in range(self.size):
        for j in range(self.size):
            self.state = (i,j)
            if self.state == self.goal:
                policy[i][j] = -1
                continue
            else:
                nbrs = self.getNeighbors()
                values = []
                for n in nbrs:
                    if n != None:
                        q, r = n[0], n[1]
                        values.append(V[q][r])
                    else:
                        values.append(-1)
                policy[i][j] = np.argmax(values)
    return policy

```